



UNC CHARLOTTE

*The* WILLIAM STATES LEE COLLEGE *of* ENGINEERING



# Real-time AI

## Lecture 1: Working with Pre-trained Networks

Hamed Tabkhi

Department of Electrical and Computer Engineering,  
University of North Carolina Charlotte (UNCC)

[htabkhiv@uncc.edu](mailto:htabkhiv@uncc.edu)

# Summary

Deep learning models automatically learn to associate inputs and desired outputs from examples.

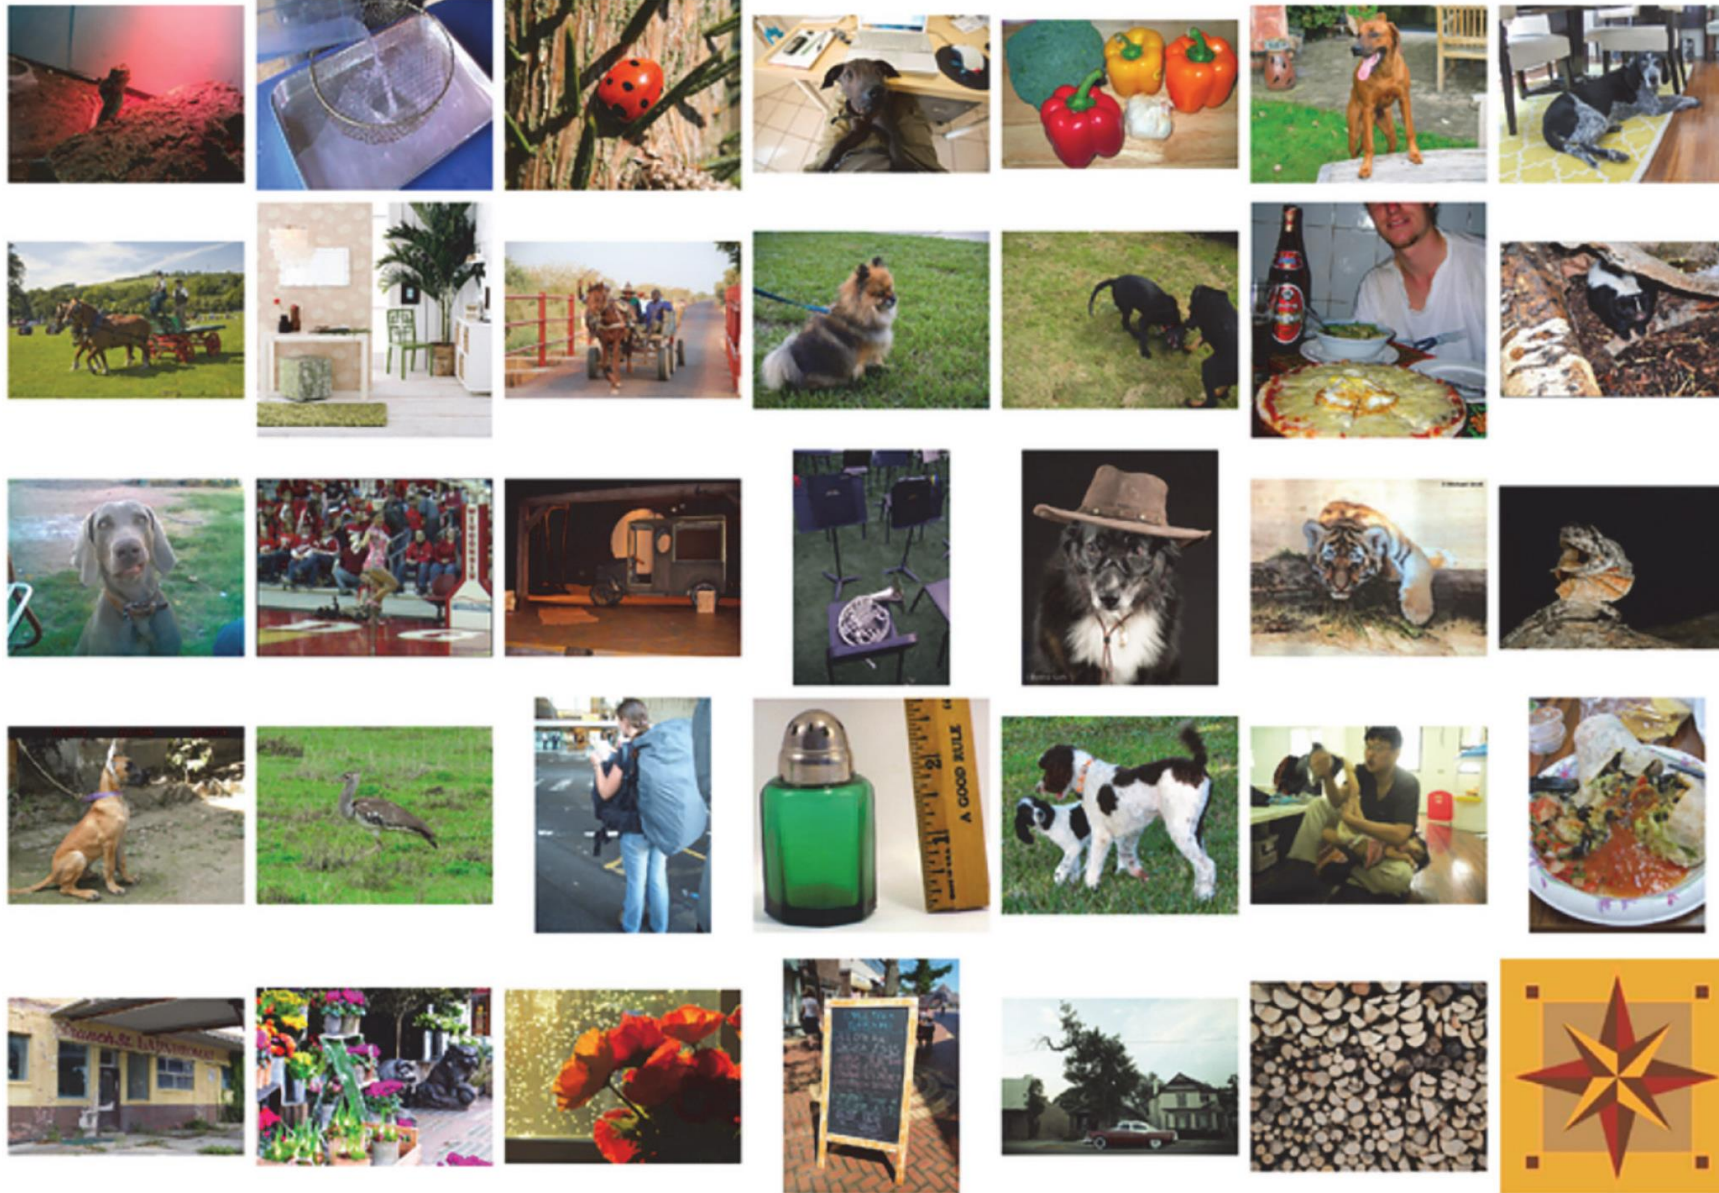
Libraries like PyTorch allow you to build and train neural network models efficiently. PyTorch minimizes cognitive overhead while focusing on flexibility and speed. It also defaults to immediate execution for operations.

TorchScript allows us to precompile models and invoke them not only from Python but also from C++ programs and on mobile devices.

Since the release of PyTorch in early 2017, the deep learning tooling ecosystem has consolidated significantly.

PyTorch provides a number of utility libraries to facilitate deep learning projects.






The WILLIAM STATES LEE COLLEGE of ENGINEERING  
UNC CHARLOTTE


# Image Classification

*A pretrained network that recognizes the subject of an image*





# AlexNet: First Image Classifier




The WILLIAM STATES LEE COLLEGE of ENGINEERING  
UNC CHARLOTTE

# AlexNet Classification Example

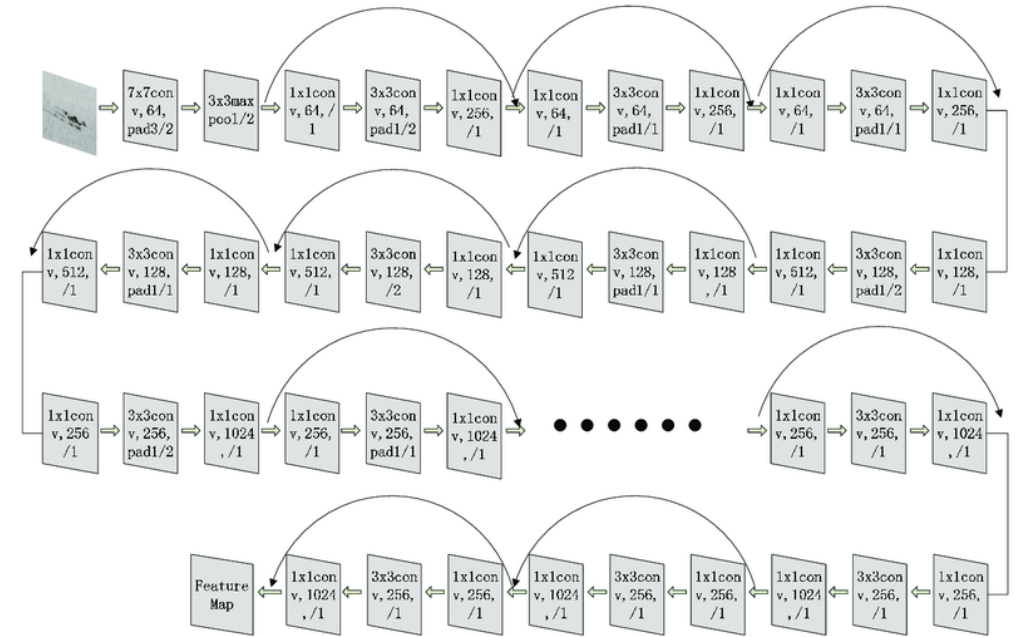


# ResNet



# ResNet 101

- A 101-layer convolutional neural network.
- Residual networks pulled a trick that made it possible to train a network composing out of 101 layers.
- resnet101 sports 44.5 million parameters—that's a lot of parameters to optimize automatically!
- It beat several benchmarks.





# Time to Try Some Codes

## **Personal Computer:**

Install Pytorch with respect to your target OS. Instruction here:

<https://pytorch.org/get-started/locally/>

Install Jupyter Notebook or Jupyter Lab

<https://jupyter.org/install>

## **Google Co-lab:**

Everything has already installed. Instruction to work with google Co-lab provided here:

[https://course19.fast.ai/start\\_colab.html](https://course19.fast.ai/start_colab.html)



*The* WILLIAM STATES LEE COLLEGE *of* ENGINEERING  
UNC CHARLOTTE

# Creating ResNet 101 Instance

Let's create an instance of the network now.

Download all source codes from

<https://github.com/deep-learning-with-pytorch/dlwpt-code>

Code related to ResNet 101 are in:

[https://github.com/deep-learning-with-pytorch/dlwpt-code/blob/master/p1ch2/2\\_pre\\_trained\\_networks.ipynb](https://github.com/deep-learning-with-pytorch/dlwpt-code/blob/master/p1ch2/2_pre_trained_networks.ipynb)

# In[4]:

```
resnet = models.resnet101(pretrained=True)
```

We'll pass an argument that will instruct the function to download the weights of `resnet101` trained on the ImageNet dataset, with 1.2 million images and 1,000 categories:



*The* WILLIAM STATES LEE COLLEGE *of* ENGINEERING  
UNC CHARLOTTE

# Creating ResNet 101 Instance

```
# In[5]: resnet

# Out[5]:
ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
    bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
    track_running_stats=True)
  (relu): ReLU(inplace)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
    ceil_mode=False)
  (layer1): Sequential(
    (0): Bottleneck(
  ...
    )
  )
  (avgpool): AvgPool2d(kernel_size=7, stride=1, padding=0)
  (fc): Linear(in_features=2048, out_features=1000, bias=True)
)
```

**modules** one per line: they are individual operations, the building blocks of a neural network. They are also called *layers* in other deep learning frameworks.

**Bottleneck Module:** group of layer modules

Sequential cascade of filters and nonlinear functions, ending with a layer (`fc`) producing scores for each of the 1,000 output classes (`out_features`). See full ResNet 101 code here:

[https://github.com/deep-learning-with-pytorch/dlwpt-code/blob/master/p1ch2/2\\_pre\\_trained\\_networks.ipynb](https://github.com/deep-learning-with-pytorch/dlwpt-code/blob/master/p1ch2/2_pre_trained_networks.ipynb)



The WILLIAM STATES LEE COLLEGE of ENGINEERING  
UNC CHARLOTTE

# Creating ResNet 101 Instance

we have to preprocess the input images so they are the right size and so that their values (colors) sit roughly in the same numerical range.

```
# In[6]:  
from torchvision import transforms  
preprocess =  
transforms.Compose([  
    transforms.Resize(256),  
    transforms.CenterCrop(224),  
    transforms.ToTensor(),  
    transforms.Normalize(  
        mean=[0.485, 0.456, 0.406],  
        std=[0.229, 0.224, 0.225]  
    )])
```

The `torchvision` module provides `transforms`, which allow us to quickly define pipelines of basic preprocessing functions:

We scale the input image to  $256 \times 256$ , crop the image to  $224 \times 224$  around the center,

It transform image to a tensor (a PyTorch multidimensional array: A 3D array with color, height, and width)

It also normalizes its RGB (red, green, blue) components so that they have defined means and standard deviations

**Note:** These need to match what was presented to the network during training



The WILLIAM STATES LEE COLLEGE of ENGINEERING  
UNC CHARLOTTE



# Creating ResNet 101 Instance

Grab a picture of our favorite dog (say, bobby.jpg from the GitHub repo).

We can start by loading an image from the local filesystem using Pillow (<https://pillow.readthedocs.io/en/stable>)

An image-manipulation module for Python:

```
# In[7]:  
from PIL import Image  
img = Image.open("../data/p1ch2/bobby.jpg")
```

```
# In[8]:  
img
```

The `resnet` variable can be called like a function, taking as input one or more images and producing an equal number of scores for each of the 1,000 ImageNet classes.



*The* WILLIAM STATES LEE COLLEGE *of* ENGINEERING  
UNC CHARLOTTE

# Creating ResNet 101 Instance

Next, we can pass the image through our preprocessing pipeline:

```
# In[9]:  
img_t = preprocess(img)
```

Then we can reshape, crop, and normalize the input tensor in a way that the network expects.

We'll understand more of this in the next two chapters; hold tight for now:

```
# In[10]:  
import torch  
batch_t = torch.unsqueeze(img_t, 0)
```

The process of running a trained model on new data is called ***inference*** in deep learning circles. In order to do inference, we need to put the network in `eval` mode:

```
# In[11]:  
resnet.eval()
```



# Creating ResNet 101 Instance

Now, we run the network

```
# In[12]:  
out = resnet(batch_t)  
out  
# Out[12]:  
tensor([[ -3.4803, -1.6618, -2.4515, -3.2662, -3.2466, -1.3611,  
-2.0465, -2.5112, -1.3043, -2.8900, -1.6862, -1.3055,  
...  
2.8674, -3.7442, 1.5085, -3.2500, -2.4894, -0.3354,  
0.1286, -1.1355, 3.3969, 4.4584]])
```

A staggering set of operations involving 44.5 million parameters has just happened, producing a vector of 1,000 scores, one per ImageNet class.



*The* WILLIAM STATES LEE COLLEGE *of* ENGINEERING  
UNC CHARLOTTE

# Creating ResNet 101 Instance

We now need to find out the label of the class that received the highest score.

To see the list of predicted labels, we load a text file listing the labels in the same order they were presented to the network during training:

```
# In[13]:  
with open('../data/p1ch2/imagenet_classes.txt') as f:  
    labels = [line.strip() for line in f.readlines()]
```

We can do that using the `max` function in PyTorch, which outputs the maximum value in a tensor as well as the indices where that maximum value occurred:

```
# In[14]:  
_, index = torch.max(out, 1)
```





# Creating ResNet 101 Instance

We also use `torch.nn.functional.softmax` (<http://mng.bz/BYnq>) to normalize our outputs to the range [0, 1], and divide by the sum.

That gives us something roughly akin to the confidence that the model has in its prediction.

```
percentage =  
torch.nn.functional.softmax(out, dim=1)[0]  
* 100  
labels[index[0]],  
percentage[index[0]].item()  
# Out[15]:  
('golden retriever', 96.29334259033203)
```



In this case, the model is 96% certain that it knows what it's looking at is a golden retriever.



*The* WILLIAM STATES LEE COLLEGE *of* ENGINEERING  
UNC CHARLOTTE

# Creating ResNet 101 Instance

Since the model produced scores, we can also find out what the second best, third best, and so on were. To do this, we can use the `sort` function, which sorts the values in ascending or descending order and also provides the indices of the sorted values in the original array:

```
# In[16]:
_, indices = torch.sort(out, descending=True)
[(labels[idx], percentage[idx].item()) for idx in indices[0][:5]]
# Out[16]:
[('golden retriever', 96.29334259033203),
 ('Labrador retriever', 2.80812406539917),
 ('cocker spaniel, English cocker spaniel, cocker', 0.28267428278923035),
 ('redbone', 0.2086310237646103),
 ('tennis ball', 0.11621569097042084)]
```

We see that the first four are dogs!

